Timed Mode Set 1 – AZ-104 Azure Administrator

YOUR TIME:
00:17:25

You have reached 54 of 79 point(s), (68.35%)

Categories

AZ-104 – Deploy and Manage Azure Compute Resources	73.68%
AZ-104 – Implement and Manage Storage	68.18%
AZ-104 – Implement and Manage Virtual Networking	62.5%
AZ-104 – Manage Azure Identities and Governance	57.14%
AZ-104 – Monitor and Maintain Azure Resources	85.71%

Sorry, you failed the test. Carefully read our detailed explanations including the references and cheat sheets then try again. 🙂

To view your record of all previous attempts:

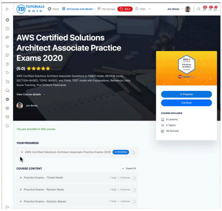

AWS Certified Solutions Architect Associate Practice Exams 2020

(5.0) ★★★★★

AWS Certified Solutions Architect Associate Questions in TIMED mode, REVIEW mode, SECTION-BASED, TOPIC-BASED, and FINAL TEST mode with Explanations, Reference Links, Score Tracking, Plus BONUS Flashcards

View Course details

COURSE INCLUDES
- 6 Lessons
- 5 Topics
- 30 Quizzes

You are enrolled in this course

YOUR PROGRESS
- AWS Certified Solutions Architect Associate Practice Exams 2020

COURSE CONTENT
- Practice Exams - Timed Mode 1 Topic | 6 Quizzes
- Practice Exams - Review Mode 1 Topic | 6 Quizzes
- Practice Exams - Section-Based 1 Topic | 4 Quizzes

Visit our FAQ page for more information on the site's features.

View Incorrect Answers

Restart Quiz

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26
27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52

✓ Answered ⚑ For Review

3. QUESTION

Your company has a Microsoft Entra ID tenant named tutorialsdojo.onmicrosoft.com and a public DNS zone for tutorialsdojo.com.

You added the custom domain name tutorialsdojo.com to Microsoft Entra ID. You need to verify that Azure can verify the domain name.

What DNS record type should you use?

● RRSIG

○ A

○ SOA

○ MX

Incorrect

Microsoft Entra ID is a cloud-based identity and access management service that enables your employees access external resources. Example resources include Microsoft 365, the Azure portal, and thousands of other SaaS applications.

Microsoft Entra ID also helps them access internal resources like apps on your corporate intranet, and any cloud apps developed for your own organization.

Home > Default Directory >

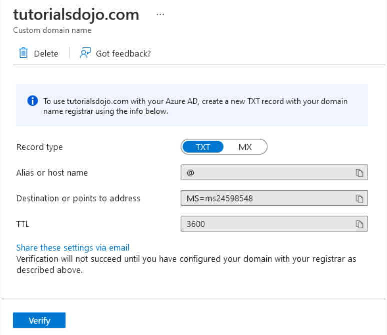

tutorialsdojo.com ···

Custom domain name

To use tutorialsdojo.com with your Azure AD, create a new TXT record with your domain name registrar using the info below.

Record type: TXT | MX

Alias or host name: @

Destination or points to address: MS=ms24598548

TTL: 3600

Share these settings via email

Verification will not succeed until you have configured your domain with your registrar as described above.

Every new Microsoft Entra ID tenant comes with an initial domain name, *<domainname>.onmicrosoft.com*. You can't change or delete the initial domain name, but you can add your organization's names. Adding custom domain names helps you to create user names that are familiar to your users, such as *azure@tutorialsdojo.com*.

You can verify your custom domain name by using TXT or MX record types.

Hence, the correct answer is: **MX.**

A is incorrect. A records are used to map domain names to IP addresses and are unrelated to domain verification in Microsoft Entra ID.

RRSIG is incorrect. RRSIG records are used in DNSSEC (Domain Name System Security Extensions) to provide cryptographic signatures for DNS records. These signaturess simply validate the authenticity of DNS data and not used for ownership verification.

SOA is incorrect. SOA records provide administrative details about the domain. This record is not relevant for domain verification.

References:

https://learn.microsoft.com/en-us/entra/fundamentals/whatis

https://learn.microsoft.com/en-us/entra/fundamentals/add-custom-domain

Check out this Microsoft Entra ID Cheat Sheet:

https://tutorialsdojo.com/microsoft-entra-id/

4. QUESTION

You have an existing Azure subscription that has the following Azure Storage accounts.

Name	Kind	Performance	Replication	Access Tier
tdaccount1	General-purpose V2	Standard	Locally-redundant storage (LRS)	Cool
tdaccount2	General-purpose V2	Premium	Read-access geo-redundant storage (RA-GRS)	Hot
tdaccount3	General-purpose V1	Premium	Geo-redundant storage (GRS)	None
tdaccount4	BlobStorage	Standard	Locally-redundant storage (LRS)	Hot

There is a requirement to identify the storage accounts that can be converted to zone-redundant storage (ZRS) replication. This must be completed through an in-place conversion without manually copying the data to another storage account.

Which of the following accounts can you convert to ZRS?

○ tdaccount1

○ tdaccount2

○ tdaccount4

● tdaccount3

Incorrect

Azure Storage always stores multiple copies of your data so that it is protected from planned and unplanned events, including transient hardware failures, network or power outages, and massive natural disasters. Redundancy ensures that your storage account meets its availability and durability targets even in the face of failures.

When deciding which redundancy option is best for your scenario, consider the tradeoffs between lower costs and higher availability. The factors that help determine which redundancy option you should choose to include are:

– How your data is replicated in the primary region.

– Whether your data is replicated to a second region that is geographically distant to the primary region, to protect against regional disasters.

– Whether your application requires read access to the replicated data in the secondary region if the primary region becomes unavailable for any reason.

Data in an Azure Storage account is always replicated three times in the primary region. Azure Storage offers four options for how your data is replicated:

1. **Locally redundant storage (LRS)** copies your data synchronously three times within a single physical location in the primary region. LRS is the least expensive replication option but is not recommended for applications requiring high availability.
2. **Zone-redundant storage (ZRS)** copies your data synchronously across three Azure availability zones in the primary region. For applications requiring high availability.
3. **Geo-redundant storage (GRS)** copies your data synchronously three times within a single physical location in the primary region using LRS. It then copies your data asynchronously to a single physical location in a secondary region that is hundreds of miles away from the primary region.
4. **Geo-zone-redundant storage (GZRS)** copies your data synchronously across three Azure availability zones in the primary region using ZRS. It then copies your data asynchronously to a single physical location in the secondary region.

You can switch a storage account from one type of replication to any other type, but some scenarios are more straightforward than others. If you want to add or remove geo-replication or read access to the secondary region, you can use the Azure portal, PowerShell, or Azure CLI to update the replication setting. However, if you want to change how data is replicated in the primary region, by moving from LRS to ZRS or vice versa, then you must perform a manual migration.

The following table provides an overview of how to switch from each type of replication to another:

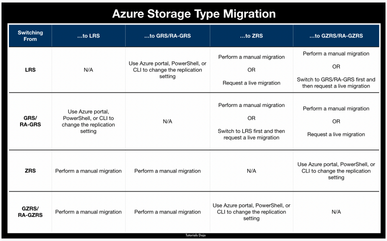

Azure Storage Type Migration

Switching From	...to LRS	...to GRS/RA-GRS	...to ZRS	...to GZRS/RA-GZRS
LRS	N/A	Use Azure portal, PowerShell, or CLI to change the replication setting	Perform a manual migration OR Request a live migration	Perform a manual migration OR Switch to GRS/RA-GRS first and then request a live migration
GRS/ RA-GRS	Use Azure portal, PowerShell, or CLI to change the replication setting	N/A	Perform a manual migration OR Switch to LRS first and then request a live migration	Perform a manual migration OR Request a live migration
ZRS	Perform a manual migration	Perform a manual migration	N/A	Use Azure portal, PowerShell, or CLI to change the replication setting
GZRS/ RA-GZRS	Perform a manual migration	Perform a manual migration	Use Azure portal, PowerShell, or CLI to change the replication setting	N/A

Tutorials Dojo

To request a live migration to ZRS, GZRS, or RA-GZRS, you need to migrate your storage account from LRS to ZRS in the primary region with no application downtime. To migrate from LRS to GZRS or RA-GZRS, first switch to GRS or RA-GRS and then request a live migration. Similarly, you can request a live migration from GRS or RA-GRS to GZRS or RA-GZRS. To migrate from GRS or RA-GRS to ZRS, first switch to LRS, then request a live migration.

Live migration is supported only for storage accounts that use LRS or GRS replication. If your account uses RA-GRS then you need to first change your account's replication type to either LRS or GRS before proceeding. This intermediary step removes the secondary read-only endpoint provided by RA-GRS before migration.

Hence, the correct answer is: **tdaccount1.**

tdaccount2 is incorrect because you need to first change your account's replication type to either LRS or GRS before you change to zone-redundant storage (ZRS). The requirement states that you must only request live migration.

tdaccount3 is incorrect because a general-purpose V1 storage account type does not support zone-redundant storage (ZRS) as its replication option. Only General-purpose V2, FileStorage, and BlockBlobStorage support ZRS.

tdaccount4 is incorrect because a BlobStorage account type does not support zone-redundant storage (ZRS) as its replication option. Only General-purpose V2, FileStorage, and BlockBlobStorage support ZRS.

References:

https://docs.microsoft.com/en-us/azure/storage/common/storage-redundancy

https://docs.microsoft.com/en-us/azure/storage/common/redundancy-migration

Check out these Azure Cheat Sheets:

https://tutorialsdojo.com/azure-storage-overview/

https://tutorialsdojo.com/locally-redundant-storage-lrs-vs-zone-redundant-storage-zrs/

5. QUESTION

You have two virtual networks named `TDVnet1` and `TDVnet2`. Your site-to-site VPN uses an Azure VPN Gateway named `TDGW1` with static routing to connect the on-premises network to `TDVnet1`.

Your Windows 10 computer named `TD1` has a point-to-site VPN connection to `TDVnet1`. Your virtual network peering between `TDVnet1` and `TDVnet2` was recently established.

Your tests confirm connectivity to `TDVnet2` from the on-premises network and to `TDVnet1` from `TD1`. Your computer `TD1` is currently unable to access `TDVnet2`.

Which action is required to enable a connection from `TD1` to `TDVnet2` ?

- ○ Restart `TDGW1` to re-establish the connection.
- ○ Enable transit gateway for `TDVnet1` .
- ● Enable transit gateway for `TDVnet2` .
- ○ Download the VPN client configuration package and re-install it on `TD1` .

Incorrect

Point-to-Site (P2S) VPN connection allows you to create a secure connection to your virtual network from an individual client computer. A P2S connection is established by starting it from the client's computer. This solution is useful for telecommuters who want to connect to Azure VNets from a remote location, such as from home or a conference. P2S VPN is also a helpful solution to utilize instead of S2S VPN when you have only a few clients that need to connect to a VNet.

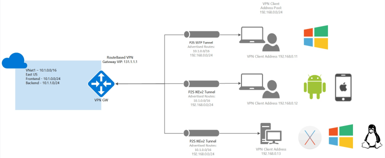

As part of the Point-to-Site configuration, you install a certificate and a VPN client configuration package which are contained in a zip file. Configuration files provide the settings required for a native Windows, Mac IKEv2 VPN, or Linux clients to connect to a virtual network over Point-to-Site connections that use native Azure certificate authentication and are specific to the VPN configuration for the virtual network.

Take note that after creating the point-to-site connection between `TD1` and `TDVnet1` , there is already a change in network topology when you created the virtual network peering with `TDVnet1` and `TDVnet2` . Whenever there is a change in the topology of your network, you will always need to download and re-install the VPN configuration file.

Hence, the correct answer is: **Download the VPN client configuration package and re-install it on `TD1` .**

The option that says: **Restart `TDGW1` to re-establish the connection** is incorrect because restarting the VPN gateway is only done when you lose cross-premises VPN connectivity on one or more Site-to-Site VPN tunnels. In this scenario, `TD1` can connect to `TDVnet1` which implies that `TDGW1` is working and running.

The options that say: **Enable transit gateway for `TDVnet1`** and **Enable transit gateway for `TDVnet2`** are incorrect. Transit gateway is a peering property that lets one virtual network use the VPN gateway in the peered virtual network for cross-premises or VNet-to-VNet connectivity. Since `TDVnet2` can connect to the on-premises network, it means that the transit gateway is already enabled and as such, enabling the transit gateway is not necessary.

References:

https://azure.microsoft.com/en-us/services/vpn-gateway/

https://docs.microsoft.com/en-us/azure/vpn-gateway/point-to-site-about

Check out this Azure VPN Gateway Cheat Sheet:

https://tutorialsdojo.com/azure-vpn-gateway/

6. QUESTION

You have a file share in your Azure subscription named `Manila-Subscription-01` .

You plan to sync files from your on-premises `TDFileServer1` to Azure.

You created an Azure file share and a storage sync service to support future Azure Backup Instant Restore capabilities.

Which four actions should you perform in sequence to synchronize files from `TDFileServer1` to Azure?

Instructions: To answer, drag the appropriate item from the column on the left to its description on the right. Each correct match is worth one point.

ANSWER OPTIONS

ANSWER AREA

Register `TDFileServer1` with Storage Sync Service	4
Deploy the Azure File Sync agent to `TDFileServer1`	3
Create a server endpoint	1

Set up a sync group and configure a cloud endpoint.	2

Incorrect

0 / 4 Points

Azure Files enables you to set up highly available network file shares that can be accessed by using the standard Server Message Block (SMB) protocol. That means that multiple VMs can share the same files with both read and write access. You can also read the files using the REST interface or the storage client libraries.

One thing that distinguishes Azure Files from files on a corporate file share is that you can access the files from anywhere in the world using a URL that points to the file and includes a shared access signature (SAS) token. You can generate SAS tokens; they allow specific access to a private asset for a specific amount of time.

File shares can be used for many common scenarios:

1. Many on-premises applications use file shares. This feature makes it easier to migrate those applications that share data to Azure. If you mount the file share to the same drive letter that the on-premises application uses, the part of your application that accesses the file share should work with minimal, if any, changes.

2. Configuration files can be stored on a file share and accessed from multiple VMs. Tools and utilities used by multiple developers in a group can be stored on a file share, ensuring that everybody can find them and that they use the same version.

3. Resource logs, metrics, and crash dumps are just three examples of data that can be written to a file share and processed or analyzed later.

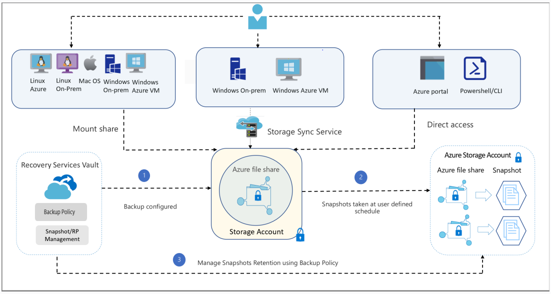

You can use Azure File Sync to centralize your organization's file shares in Azure Files while keeping the flexibility, performance, and compatibility of an on-premises file server. Azure File Sync transforms Windows Server into a quick cache of your Azure file share. You can use any protocol that's available on Windows Server to access your data locally, including SMB, NFS, and FTPS. You can have as many caches as you need across the world.

You can sync `TDFileServer1` to Azure using the following steps in order:

1. Prepare Windows Server to use with Azure File Sync

   – You need to disable Internet Explorer Enhanced Security Configuration in your server. This is required only for initial server registration. You can re-enable it after the server has been registered.

2. Deploy the Storage Sync Service

   – Allows you to create sync groups that contain Azure file shares across multiple storage accounts and multiple registered Windows Servers.

3. Deploy the Azure File Sync agent to `TDFileServer1`

   – The Azure File Sync agent is a downloadable package that enables Windows Server to be synced with an Azure file share.

4. Register `TDFileServer1` with Storage Sync Service

   – This establishes a trust relationship between your server (or cluster) and the Storage Sync Service. A server can only be registered to one Storage Sync Service and can sync with other servers and Azure file shares associated with the same Storage Sync Service.

5. Set up a sync group and configure a cloud endpoint.

   – A sync group defines the sync topology for a set of files. Endpoints within a sync group are kept in sync with each other.

6. Create a server endpoint

   – A server endpoint represents a specific location on a registered server, such as a folder on a server volume.

Hence, the correct order of deployment is:

1. **Deploy the Azure File Sync agent to `TDFileServer1`.**
2. **Register `TDFileServer1` with Storage Sync Service.**
3. **Set up a sync group and configure a cloud endpoint.**
4. **Create a server endpoint.**

References:

https://docs.microsoft.com/en-us/azure/storage/files/storage-files-introduction

https://docs.microsoft.com/en-us/azure/storage/files/storage-sync-files-deployment-guide

Check out this Azure Files Cheat Sheet:

https://tutorialsdojo.com/azure-file-storage/

13. QUESTION

You have a server in your on-premises datacenter that contains a DNS server named `TD1` with a primary DNS zone for the `tutorialsdojo.com` domain.

You have an Azure subscription named `TD-Subscription1` .

You plan to migrate the `tutorialsdojo.com` zone to an Azure DNS zone in `TD-Subscription1` . You must ensure that you minimize administrative effort.

Which tools can you use? (Select TWO.)

☐ Azure Portal

☐ Azure Resource Manager templates

☑ Azure CLI

☑ Azure PowerShell

☐ Azure CloudShell

Incorrect

Azure DNS is a hosting service for DNS domains that provides name resolution by using Microsoft Azure infrastructure. By hosting your domains in Azure, you can manage your DNS records by using the same credentials, APIs, tools, and billing as your other Azure services.

You can't use Azure DNS to buy a domain name. For an annual fee, you can buy a domain name by using App Service domains or a third-party domain name registrar. Your domains can then be hosted in Azure DNS for record management.

```
c:\>az login
c:\>az account set --subscription xxxxxxxx-yyyy-aaaa-rrrr-nnnnnnnnnnnn
C:\>az network dns zone import -g ResourceGroup-n mydomain.gr -f "C:\dns\mydomainexport.txt"
== BEGINNING ZONE IMPORT: mydomain.gr==

(1/6) Imported 1 records of type 'soa' and name '@'
(3/6) Imported 2 records of type 'NS' and name '@'
(4/6) Imported 1 records of type 'a' and name '@'
(5/6) Imported 1 records of type 'a' and name 'g'
(6/6) Imported 1 records of type 'a' and name 'www'

== 6/6 RECORDS IMPORTED SUCCESSFULLY: 'mydomain.gr' ==
```

A DNS zone file is a text file that contains details of every Domain Name System (DNS) record in the zone. It follows a standard format, making it suitable for transferring DNS records between DNS systems. Using a zone file is a quick, reliable, and convenient way to transfer a DNS zone into or out of Azure DNS.

Azure DNS supports importing and exporting zone files using the **Azure CLI** and the **Azure portal**. Zone file import and export are **not** available in Azure PowerShell.

Note: Older guidance (and many third-party study resources) list the Azure CLI as the *only* method for zone file import. Microsoft has since added native import/export support in the Azure portal, so both the Azure CLI and the Azure portal are now valid tools.

Hence, the correct answers are:

– Azure CLI

– Azure Portal

Azure PowerShell is incorrect because it has no native cmdlet for importing a DNS zone file. It can manage individual DNS records and zones, but bulk zone file import via PowerShell is not supported by Azure DNS.

Azure Resource Manager templates is incorrect because ARM templates are primarily used to declaratively deploy and configure Azure infrastructure, not to migrate or import data. There is no ARM template mechanism to ingest a BIND-format DNS zone file.

Azure CloudShell is incorrect because it is an execution environment rather than a separate import method. You can run the `az network dns zone import` command inside Cloud Shell, but the import is performed by the Azure CLI itself. The two tools that natively provide zone file import are the Azure CLI and the Azure portal.

References:

https://docs.microsoft.com/en-us/azure/dns/dns-overview

https://docs.microsoft.com/en-us/azure/dns/dns-import-export

Check out these Azure DNS and Azure User Tools Cheat Sheets:

https://tutorialsdojo.com/azure-dns/

https://tutorialsdojo.com/azure-user-tools/

14. QUESTION

You have an Azure subscription that contains the following virtual network peerings:

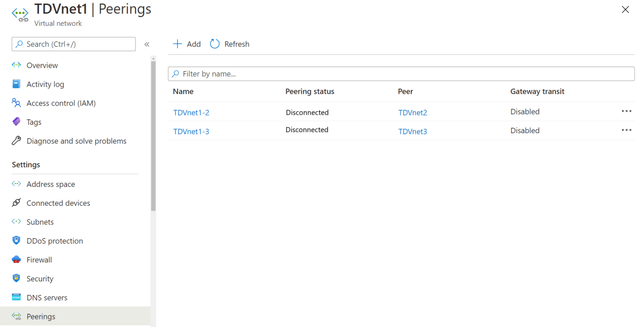

Name	Peering status	Peer	Gateway transit
TDVnet1-2	Disconnected	TDVnet2	Disabled
TDVnet1-3	Disconnected	TDVnet3	Disabled

Select the correct answer from the drop-down list of options. Each correct selection is worth one point.

Virtual Machines on TDVnet1 has network connectivity with hosts on: TDVnet1, TDVnet2, and TDVnet3 (TDVnet1 only)

What is the first thing you need to do to change the status of the peering connection for TDVnet2 to `Connected` :

Enable Allow gateway transit setting (Delete TDVnet1-2)

Incorrect

0 / 2 Points

Azure Virtual Network (VNet) is the fundamental building block for your private network in Azure. VNet enables many types of Azure resources, such as Azure Virtual Machines (VM), to securely communicate with each other, the Internet, and on-premises networks. VNet is similar to a traditional network that you'd operate in your own datacenter but brings with it additional benefits of Azure's infrastructure such as scale, availability, and isolation.

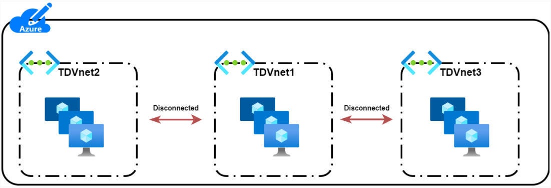

Virtual network peering enables you to connect two or more Virtual Networks in Azure seamlessly. The virtual networks appear as one for connectivity purposes. The traffic between virtual machines in peered virtual networks uses the Microsoft backbone infrastructure. Like traffic between virtual machines in the same network, traffic is routed only through Microsoft's private network.

In the image above, `TDVnet1` is the hub while `TDVnet2` and `TDVnet3` are the spoke. `TDVnet1` hosts can not communicate with `TDVnet2` and `TDVnet3` because their peerings are in a disconnected state.

Take note that if your VNet peering connection is in a Disconnected state, it means one of the links created was deleted. To re-establish a peering connection, you will need to delete the disconnected peer and recreate it.

Therefore, virtual machines on `TDVnet1` can communicate to hosts on **TDVnet1 only** because the peerings associated with `TDVnet1` are in a disconnected state. It means that traffic between virtual networks is prohibited.

Conversely, you need to **Delete TDVnet1-2** to re-establish the connection. Once you have deleted the disconnected peer, you can then recreate it.

Hence, the correct answers are:

– Virtual Machines on TDVnet1 has network connectivity with hosts on: TDVnet1 only

– What is the first thing you need to do to change the status of the peering connection for TDVnet2 to `Connected` : Delete TDVnet1-2

The following options are incorrect because `TDVnet2` and `TDVnet3` have a disconnected peer with `TDVnet1` . No traffic will be able to flow between virtual networks as long as the peer's status is disconnected. To re-establish the connection, you must delete the disconnected peer and recreate it.

– TDVnet2 only

– TDVnet3 only

– TDVnet1,TDVnet2, and TDVnet3

The option that says: **Change the address space** is incorrect because you can not typically change the address space of a virtual network if there is an active peering connection. You need to delete the peer first to change the address space.

The option that says: **Delete a subnet** is incorrect because even if you delete or add a subnet, it will not have any impact on the state of the peering connection.

The option that says: **Enable Allow gateway transit setting** is incorrect because this feature is simply a peering property that lets one virtual network use the

VPN gateway in the peered virtual network for cross-premises or VNet-to-VNet connectivity.

References:

https://docs.microsoft.com/en-us/azure/virtual-network/virtual-networks-overview

https://docs.microsoft.com/en-us/azure/virtual-network/virtual-network-peering-overview

Check out this Azure Virtual Network Cheat Sheet:

https://tutorialsdojo.com/azure-virtual-network-vnet/

15. QUESTION

Your company has an Azure subscription that contains a virtual machine named `TD1` .

You need to connect to `TD1` from a computer connected to the Internet named `Workstation1` . `TD1` has a network security group with the following inbound security rules:

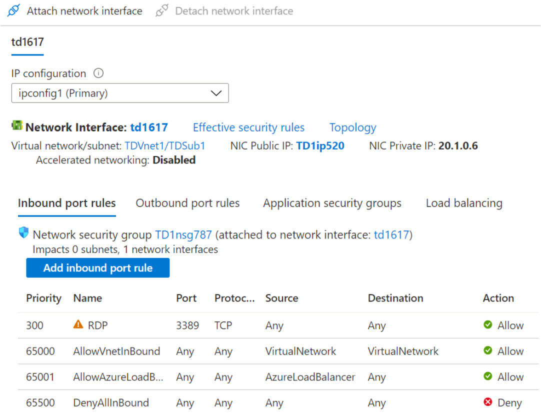

IP configuration: ipconfig1 (Primary)

Network Interface: td1617

Virtual network/subnet: TDVnet1/TDSub1 NIC Public IP: **TD1ip520** NIC Private IP: **20.1.0.6**
Accelerated networking: **Disabled**

Network security group TD1nsg787 (attached to network interface: td1617)
Impacts 0 subnets, 1 network interfaces

Priority	Name	Port	Protoc...	Source	Destination	Action
300	⚠ RDP	3389	TCP	Any	Any	✓ Allow
65000	AllowVnetInBound	Any	Any	VirtualNetwork	VirtualNetwork	✓ Allow
65001	AllowAzureLoadB...	Any	Any	AzureLoadBalancer	Any	✓ Allow
65500	DenyAllInBound	Any	Any	Any	Any	✗ Deny

You try to connect to `TD1` using `Workstation1` but you are unable to.

What should you do first to establish a Remote Desktop connection to `TD1` ?

○ Modify the priority of the RDP rule

○ Remove the RDP rule

● Redeploy `TD1`

○ Start `TD1`

Incorrect

Azure Virtual Network (VNet) is the fundamental building block for your private network in Azure. VNet enables many types of Azure resources, such as Azure Virtual Machines (VM), to securely communicate with each other, the Internet, and on-premises networks. VNet is similar to a traditional network that you'd operate in your own datacenter but brings with it additional benefits of Azure's infrastructure such as scale, availability, and isolation.

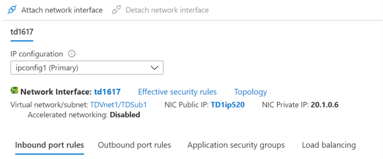

IP configuration: ipconfig1 (Primary)

Network Interface: td1617

Virtual network/subnet: TDVnet1/TDSub1 NIC Public IP: **TD1ip520** NIC Private IP: **20.1.0.6**
Accelerated networking: **Disabled**

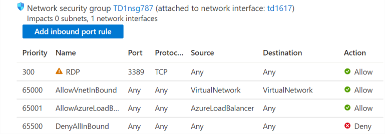

Network security group TD1nsg787 (attached to network interface: td1617)
Impacts 0 subnets, 1 network interfaces

Priority	Name	Port	Protoc...	Source	Destination	Action
300	⚠ RDP	3389	TCP	Any	Any	✓ Allow
65000	AllowVnetInBound	Any	Any	VirtualNetwork	VirtualNetwork	✓ Allow
65001	AllowAzureLoadB...	Any	Any	AzureLoadBalancer	Any	✓ Allow
65500	DenyAllInBound	Any	Any	Any	Any	✖ Deny

Public IP addresses allow Internet resources to communicate inbound to Azure resources. Public IP addresses enable Azure resources to communicate to the Internet and public-facing Azure services. The address is dedicated to the resource until it's unassigned by you. A resource without a public IP assigned can communicate outbound. Azure dynamically assigns an available IP address that isn't dedicated to the resource.

There are two types of public IP address assignment methods: **static** and **dynamic**.

Static addresses are assigned when a public IP address is created. Static addresses are not released until a public IP address resource is deleted. If the address is not associated to a resource, you can change the assignment method after the address is created. If the address is associated to a resource, you may not be able to change the assignment method.

Dynamic addresses are assigned only after a public IP address is associated to an Azure resource, and the resource is started for the first time. Dynamic addresses can change if they're assigned to a resource, such as a virtual machine, and the virtual machine is stopped (deallocated) and then started.

The address remains the same if a virtual machine is rebooted or stopped (but not deallocated). Dynamic addresses are released when a public IP address resource is dissociated from a resource it is associated with.

Take note that since `TD1` does not have an assigned IP address, it only means that `TD1` is deallocated (stopped). Remember, if you want to retain your public IP address, you need to use a static IP address.

Hence, the correct answer is: **Start `TD1`** .

The option that says: **Remove the RDP rule** is incorrect because the RDP rule already allows Remote Desktop connection from the Internet. Removing the RDP rule will prevent `Workstation1` from establishing a Remote Desktop connection to `TD1` once it is in a `Running` state.

The option that says: **Modify the priority of the RDP rule** is incorrect. Since this is the highest rule, it will be the first one that the network security group will evaluate.

The option that says: **Redeploy `TD1`** is incorrect because redeploying `TD1` will only add to the administrative effort. It would be best if you simply started `TD1` for it to receive a public IP address.

References:

https://docs.microsoft.com/en-us/azure/virtual-network/virtual-networks-overview

https://docs.microsoft.com/en-us/azure/virtual-network/public-ip-addresses

Check out this Azure Virtual Network Cheat Sheet:

https://tutorialsdojo.com/azure-virtual-network-vnet/

16. QUESTION

You have an Azure subscription named `TDSubscription` that contains Azure file share called `TDShare1` and `TDShare2` . Both have the same storage account and the same region.

You deploy the following resources:

Name	Service
TDSync1	Storage Sync Service
TDGroup1	Sync Group

`TDShare1` is the cloud endpoint of `TDGroup1`
You plan to back up the following file servers in your on-premises datacenter to Azure:

Name	Folders	Registered to
FileServer1	E:\tutorials, F:\dojo	TDSync1
FileServer2	C:\files	TDSync1

You then add `E:\tutorials` of `FileServer1` as the server endpoint of `TDGroup1` .

For each of the following items, choose Yes if the statement is true or choose No if the statement is false. Take note that each correct item is worth one point.

Questions	Yes	No
You can add `C:\files` of `FileServer2` as a server endpoint of `TDGroup1`	○	●
You can add `F:\dojo` of `FileServer1` as a server endpoint to `TDGroup1`	●	○

You can add `TDShare2` to `TDGroup1` as a cloud endpoint	○	●

Incorrect

1 / 3 Points

Azure Files enables you to set up highly available network file shares that can be accessed by using the standard Server Message Block (SMB) protocol. That means that multiple VMs can share the same files with both read and write access. You can also read the files using the REST interface or the storage client libraries.

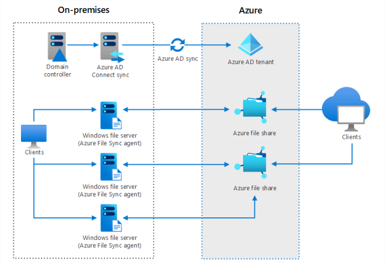
On-premises
Azure
Domain controller
Azure AD Connect sync
Azure AD sync
Azure AD tenant
Clients
Windows file server (Azure File Sync agent)
Azure file share
Clients
Windows file server (Azure File Sync agent)
Azure file share
Windows file server (Azure File Sync agent)

A sync group defines the sync topology for a set of files. Endpoints within a sync group are kept in sync with each other. A sync group must contain one cloud endpoint, which represents an Azure file share, and one or more server endpoints.

A cloud endpoint is a pointer to an Azure file share. All server endpoints will sync with a cloud endpoint, making the cloud endpoint the hub.

A server endpoint represents a specific location on a registered server, such as a folder on a server volume.

Take note that multiple server endpoints can exist on the same volume if their namespaces are not overlapping (for example, F:\sync1 and F:\sync2) and each endpoint is syncing to a unique sync group meaning you can not have more than one server endpoint from the same server in the same sync group.

The statement that says: **You can add `C:\files` of `FileServer2` as a server endpoint of `TDGroup1`** is correct because `FileServer2` has no server endpoint yet on `TDGroup1`. Therefore, you can add the file server to the sync group without any restrictions.

The statement that says: **You can add `TDShare2` to `TDGroup1` as a cloud endpoint** is incorrect because you can only have one cloud endpoint per sync group. If you want to add another cloud endpoint, you must create another sync group.

The statement that says: **You can add `F:\dojo` of `FileServer1` as a server endpoint to `TDGroup1`** is incorrect because `TDGroup1` already has a server endpoint for `FileServer1` for the folder `E:\tutorials`. Take note that you can not have more than one server endpoint from the same server in the same sync group. If you need to add the folder `F:\dojo` of `FileServer1`, you need to create another sync group.

References:

https://docs.microsoft.com/en-us/azure/storage/files/storage-files-introduction

https://docs.microsoft.com/en-us/azure/storage/files/storage-sync-files-deployment-guide

Check out this Azure Files Cheat Sheet:

https://tutorialsdojo.com/azure-file-storage/

20. QUESTION

You have an Azure subscription that has the following vCPU quotas.

Quota	Location	Usage
Standard Av2 Family vCPUs	North Central US	0 of 15
Standard DSv3 Family vCPUs	North Central US	0 of 15
Standard F Family vCPUs	South Central US	0 of 15
Total Regional vCPUs	North Central US	0 of 15

You plan to create the virtual machines listed below in the order they are listed.

Name	Size	vCPUs
VM4	A2m_v2	2
VM5	D2s_v3	2
VM6	A8_v2	8

The deployed virtual machines are shown in the table below.

Name	Size	vCPUs	Location	Status
VM1	A4_v2	4	North Central US	Running
VM2	D8s_v3	8	North Central US	Stopped (Deallocated)
VM3	F2s_v2	2	South Central US	Running

For each of the following items, choose **Yes** if the statement is true or choose **No** if the statement is false. Take note that each correct item is worth one point.

Questions	Yes	No
You can create VM6 in North Central US Region.	○	●
You can create VM5 in North Central US Region.	●	○
You can create VM4 in North Central US Region.	●	○

Incorrect

2 / 3 Points

Azure Virtual Machines (VM) is one of several types of on-demand, scalable computing resources that Azure offers. Typically, you choose a VM when you need more control over the computing environment. An Azure VM gives you the flexibility of virtualization without having to buy and maintain the physical hardware that runs it. However, you still need to maintain the VM by performing tasks such as configuring, patching, and installing the software that runs on it.

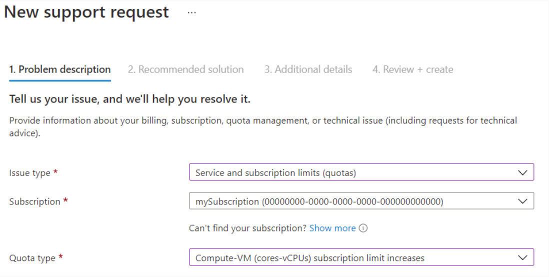

New support request ···

Tell us your issue, and we'll help you resolve it.

Provide information about your billing, subscription, quota management, or technical issue (including requests for technical advice).

Issue type * Service and subscription limits (quotas)

Subscription * mySubscription (00000000-0000-0000-0000-000000000000)

Can't find your subscription? Show more ⓘ

Quota type * Compute-VM (cores-vCPUs) subscription limit increases

The vCPU quotas for virtual machines and virtual machine scale sets are arranged in two tiers for each subscription in each region.

– Total Regional vCPUs

– VM size family cores

Every time you deploy a new VM, the vCPUs must not exceed the vCPU quota for the VM size family or the total regional vCPU quota. If either of those quotas has been exceeded, the VM deployment will not be allowed. Take note that there is also a quota for the overall number of virtual machines in the region. The quota is calculated based on the total number of cores that are currently allocated (running) in the region. If you need additional cores, you can request a quota increase or delete VMs that are no longer needed.

Hence, the correct answer is:

– You can create VM4 in North Central US Region. The remaining vCPU quota in North Central US is 11 vCPUs (15 – 4 from VM1). If you created VM4 in the North Central US Region, the total vCPUs in that Region become 6 of 15 vCPUs.

The statement that says: **You can create VM5 in the North Central US Region** is incorrect because after VM4 is created, only 1 vCPU remains in the regional quota. In the regional quota. VM5 requires 2 vCPUs, which would bring the total to 16, of 15 exceeding the regional limit.

The statement that says: **You can create VM6 in North Central US Region** is incorrect because if you create VM6 in the North Central US, it will simply exceed the total regional vCPU quota. VM6 requires 8 vCPUs, which cannot fit within the remaining quota after attempting to create the earlier VMs.

References:

https://docs.microsoft.com/en-us/azure/virtual-machines/windows/quotas

https://docs.microsoft.com/en-us/azure/virtual-machines/sizes

https://docs.microsoft.com/en-us/azure/azure-portal/supportability/per-vm-quota-requests

Check out this Azure Virtual Machines Cheat Sheet:

https://tutorialsdojo.com/azure-virtual-machines/

21. QUESTION

You plan to host a web application in three Azure virtual machines. The company already uses Azure Virtual WAN to connect its branch offices.

You need to make sure that there are at least two virtual machines running if an Azure data center becomes inaccessible.

What should you do?

○ Deploy one virtual machine in each Availability Set.

○ Deploy one virtual machine in each Availability Zone.

○ Deploy all the virtual machines in a single Availability Zone.

● Deploy all the virtual machines in a single Availability Set.

Incorrect

Azure Virtual Machines (VMs) are one of several types of on-demand, scalable computing resources that Azure offers. Typically, you choose a VM when you need more control over the computing environment. An Azure VM gives you the flexibility of virtualization without having to buy and maintain the physical hardware that runs it. However, you still need to maintain the VM by performing tasks such as configuring, patching, and installing the software that runs on it.

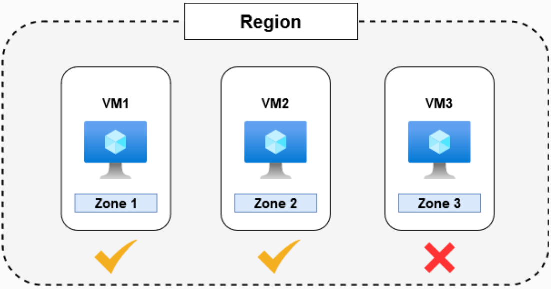

In Azure, there are two options for managing availability and resiliency for your applications. The first option is availability sets. It is used to protect applications from hardware failures within an Azure data center. Meanwhile, availability zones are used to protect applications against Azure data center failures. Take note that an availability set only protects your resources from planned and unplanned maintenance. It cannot protect your applications from data center outages. Also, in the availability set, if a hardware or software failure happens, only a subset of your VMs is impacted, and your overall solution stays operational.

For example, when you create a new VM, you specify the availability set as a parameter. Azure makes sure the VMs are isolated across multiple physical hardware resources within the data center. If the physical hardware that one of your servers is running on has a problem, you know the other instances of your servers will keep running because they're on different hardware.

Based on the given requirements, you can protect your web application from data center outages if you deploy the three virtual machines in a separate Availability Zone. Remember that Availability Zones are unique physical locations within an Azure region. Each zone is made up of one or more data centers equipped with independent power, cooling, and networking. To ensure resiliency, there is a minimum of three separate zones in all enabled regions. The physical separation of Availability Zones within a region protects applications and data from datacenter failures.

Hence, the correct answer is: **Deploy one virtual machine in each Availability Zone.**

The option that says: **Deploy all the virtual machines in a single Availability Zone** is incorrect because if the Availability Zone becomes inaccessible, then all of the resources in that location will also be affected. To achieve a highly available application, you must deploy the virtual machines in multiple Availability Zones.

The option that says: **Deploy all the virtual machines in a single Availability Set** is incorrect because an Availability Set only isolates virtual machines from each other. This means that the virtual machines are still in the same data center. To protect your application from a data center outage, you must deploy the virtual machines in three Availability Zones.

The option that says: **Deploy one virtual machine in each Availability Set** is incorrect. Deploying the virtual machines in a separate Availability Set does not mean that they are protected from a data center outage. Take note that this option only ensures that your VMs are distributed across multiple fault domains in the Azure data center. Therefore, if the data center becomes unavailable, your application becomes unavailable too.

References:

https://docs.microsoft.com/en-us/azure/virtual-machines/manage-availability#use-availability-zones-to-protect-from-datacenter-level-failures

https://docs.microsoft.com/en-us/azure/availability-zones/az-overview#availability-zones

Check out this Azure Virtual Machines Cheat Sheet:

https://tutorialsdojo.com/azure-virtual-machines/

27. QUESTION

You plan to use an Azure Resource Manager (ARM) template to deploy 5 web apps in the same region.

You are required to launch the application in the most cost-effective way.

Which of the following options fulfills this requirement?

○ Create a CDN endpoint.

○ Create one App Service plan.

● Create five App Service plans.

○ Create an Application Gateway

Incorrect

Azure Resource Manager (ARM) templates are primarily used to implement infrastructure as code for your Azure solutions. The template is a JavaScript Object Notation (JSON) file that defines your project's infrastructure and configuration. The template uses declarative syntax, which lets you state what you intend to deploy without writing the sequence of programming commands to create it. In the template, you specify the resources to deploy and the properties for those resources.

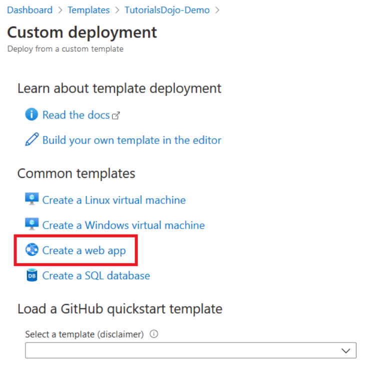

Dashboard > Templates > TutorialsDojo-Demo >

Custom deployment

Deploy from a custom template

Learn about template deployment

Read the docs

Build your own template in the editor

Common templates

Create a Linux virtual machine

Create a Windows virtual machine

Create a web app

Create a SQL database

Load a GitHub quickstart template

Select a template (disclaimer)

The main requirement in this scenario is to deploy web apps in the most cost-effective way. To accomplish this requirement, you can create one App Service plan and use the plan to deploy five web apps. If you recall the Azure App Service concepts, you can configure one or more apps to run on the same computing resources (or in the same App Service plan). Therefore, if you deploy the five web apps in the same region, you can use one App Service plan for your resources.

Hence, the correct answer is: **Create one App Service plan.**

The option that says: **Create five App Service plan** is incorrect because the requirement in this scenario is to deploy the five web apps to the same region in the most cost-effective way. This approach is applicable if you need to deploy web apps in different regions.

The option that says: **Create an Application Gateway** is incorrect because you can't deploy five web apps using Azure Application Gateway. This service is simply a web traffic load balancer and is not capable of hosting an application.

The option that says: **Create a CDN endpoint** is incorrect because a CDN endpoint only represents a specific configuration of content delivery behavior and access. You must create one App Service plan to fulfill the requirement in the scenario.

References:

https://docs.microsoft.com/en-us/azure/app-service/overview-hosting-plans

https://docs.microsoft.com/en-us/azure/azure-resource-manager/templates/overview

Check out this Azure App Service Cheat Sheet:

https://tutorialsdojo.com/azure-app-service/

31. QUESTION

Your company has a Microsoft Entra tenant named `TD-Entra-ID` that contains 3 User Administrators and 2 members of the Global Administrators group.

The administrators regularly review Account Management properties for each user.

You recently purchased 5 Premium P1 licenses.

You need to make sure that the users in your tenant have access to all the Premium P1 features.

What should you do to satisfy the above requirement?

● Configure external collaboration settings to grant access to selected users.

○ Select the user in your tenant and assign a new role in the `Directory` role blade of each user.

○ In the `Licenses` blade of Microsoft Entra ID, select the user in your tenant and assign the license.

○ Use the Bulk create user operation in Microsoft Entra ID to provision users with Premium P1 features.

Incorrect

Microsoft Entra ID is a cloud-based identity and access management service that enables your employees to access external resources. Example resources include Microsoft 365, the Azure portal, and thousands of other SaaS applications.

Microsoft Entra ID also helps them access internal resources like apps on your corporate intranet and any cloud apps developed for your own organization.

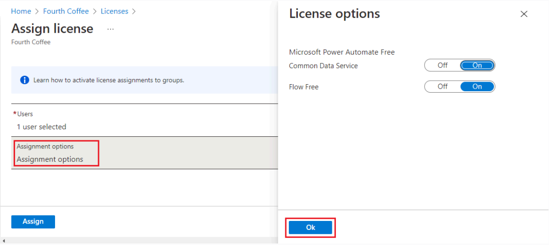

Home > Fourth Coffee > Licenses >

Assign license ...

Fourth Coffee

Learn how to activate license assignments to groups.

*Users

1 user selected

Assignment options

Assignment options

Assign

License options

Microsoft Power Automate Free

Common Data Service Off On

Flow Free Off On

Ok

There are several license plans available for the Microsoft Entra ID service, including:

– Microsoft Entra ID Free

– Microsoft Entra ID P1

– Microsoft Entra ID P2

To ensure that the users in your tenant have access to Premium P1 license features, you must manually add the license to each user or add the license to a group. Remember that only users with active licenses can access and use the licensed Microsoft Entra ID services. Also, licenses are applied per tenant, and you can't transfer them to other tenants.

Hence, the correct answer is: **In the `Licenses` blade of Microsoft Entra ID, select the user in your tenant and assign the license.**

The option that says: **Select the user in your tenant and assign a new role in the `Directory` role blade of each user** is incorrect because this approach only assigns a new role to your user. To assign the Premium P1 licenses to your users, you must go to the Licenses tab and click Assignments.

The option that says: **Use the Bulk create user operation in Microsoft Entra ID to provision users with Premium P1 features** is incorrect because bulk user creation only provisions user accounts. It does not assign licenses or activate Premium P1 features. Even if users are created in bulk, each one must still be assigned a Premium P1 license explicitly to access those features.

The option that says: **Configure external collaboration settings to grant access to selected users** is incorrect because external collaboration settings are primarily used to control how users outside your organization can access your organization's resources, such as sharing files or inviting guest users. It does not impact the internal users' access to Premium P1 features, which is managed through license assignment.

References:

https://learn.microsoft.com/en-us/entra/fundamentals/whatis

https://learn.microsoft.com/en-us/entra/fundamentals/license-users-groups

Check out this Microsoft Entra ID Cheat Sheet:

https://tutorialsdojo.com/microsoft-entra-id/

32. QUESTION

You plan to migrate your business-critical application to Azure virtual machines.

You need to make sure that at least two VMs are available during planned Azure maintenance.

What should you do?

○ Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 2 and `platformFaultDomainCount` (Fault domains) set to 3.

○ Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 3 and `platformFaultDomainCount` (Fault domains) set to 2.

○ Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 1 and `platformFaultDomainCount` (Fault domains) set to 3.

● Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 3 and `platformFaultDomainCount` (Fault domains) set to 1.

Incorrect

Azure periodically updates its platform to improve the reliability, performance, and security of the host infrastructure for virtual machines. The purpose of these updates ranges from patching software components in the hosting environment to upgrading networking components or decommissioning hardware.

Updates rarely affect the hosted VMs. When updates do have an effect, Azure chooses the least impactful method for updates:

– If the update doesn't require a reboot, the VM is paused while the host is updated, or the VM is live-migrated to an already updated host.

– If maintenance requires a reboot, you're notified of the planned maintenance. Azure also provides a time window in which you can start the maintenance yourself, at a time that works for you. The self-maintenance window is typically 35 days unless the maintenance is urgent. Azure is investing in technologies to reduce the number of cases in which planned platform maintenance requires the VMs to be rebooted.

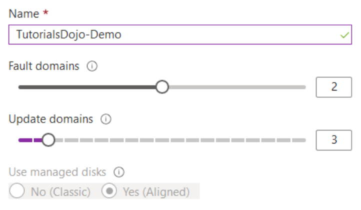

Name *
TutorialsDojo-Demo

Fault domains: 2

Update domains: 3

Use managed disks
○ No (Classic) ● Yes (Aligned)

The main objective of the question is to test your understanding of update and fault domains. Since it's a requirement in the scenario that at least two virtual machines must be available during planned maintenance, you should provision three update domains in the Availability Set. Each virtual machine in your availability set is assigned to an update domain and a fault domain.

During scheduled maintenance, only one update domain is updated at any given time. Update domains aren't necessarily updated sequentially. A rebooted update domain is given 30 minutes to recover before maintenance is initiated on a different update domain. Because planned maintenance is governed by update domains, three update domains ensure that only a portion of your VMs are rebooted at once, keeping the rest available. Fault domains protect against unplanned hardware failures, so the recommended configuration pairs multiple update domains with at least two fault domains.

You can configure these values directly using an ARM template, Bicep, and Terraform by setting the `platformUpdateDomainCount` and `platformFaultDomainCount` properties in the Availability Set resource definition. This allows you to explicitly control how your VMs are distributed across update and fault domains to meet availability requirements during planned maintenance.

Hence, the correct answer is: **Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 3 and `platformFaultDomainCount` (Fault domains) set to 2.**

The option that says: **Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 3 and `platformFaultDomainCount` (Fault domains) set to 1** is incorrect because a single fault domain leaves your VMs with no protection against rack or hardware failures. While three update domains would simply address the planned-maintenance requirement, the recommended configuration uses at least two fault domains, so 3 update domains with 2 fault domains is the better choice.

The option that says: **Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 2 and `platformFaultDomainCount` (Fault domains) set to 3** is incorrect because you should provision three update domains to keep more VMs available during planned maintenance. This option only prioritizes fault domains, which address unplanned hardware failures rather than the planned maintenance called for in the scenario.

The option that says: **Create an Availability Set with `platformUpdateDomainCount` (Update domains) set to 1 and `platformFaultDomainCount` (Fault domains) set to 3** is incorrect because a single update domain means all VMs belong to the same maintenance group and can be rebooted together during planned maintenance, leaving no VMs available. Three update domains must be provisioned to satisfy the requirement.

References:

https://learn.microsoft.com/en-us/azure/virtual-machines/maintenance-and-updates

https://learn.microsoft.com/en-us/azure/virtual-machines/availability

https://learn.microsoft.com/en-us/azure/templates/microsoft.compute/availabilitysets?pivots=deployment-language-arm-template

Check out this Azure Virtual Machines Cheat Sheet:

https://tutorialsdojo.com/azure-virtual-machines/

34. QUESTION

You manage an Azure subscription with 120 virtual machines distributed across multiple resource groups.

You aim to identify unattached disks that can be safely deleted from specific resource groups to reduce costs.

You also want to ensure you only receive recommendations relevant to those selected resource groups.

What action should you take?

● Assign the Billing Administrator role to monitor unattached disks in selected resource groups for potential deletion.

○ Use Azure Monitor to create a custom log query that identifies unattached disks in the specified resource groups.

○ Modify the Advisor configuration in Azure Advisor to include only the targeted resource groups, then review its cost optimization recommendations.

○ View the Advisor recommendations under Azure Cost Management and apply filters to focus on the resource groups in question.

Incorrect

Azure Advisor is your go-to personalized cloud consultant, dedicated to helping you optimize your Azure deployments for maximum efficiency. By thoroughly analyzing your resource configurations and monitoring usage data, Azure Advisor delivers tailored recommendations that enhance various aspects of your cloud environment. These suggestions focus on improving cost-effectiveness, boosting performance, ensuring reliability, and fortifying security, allowing you to maximize the value of your Azure resources while minimizing potential risks. With Azure Advisor, you can confidently navigate your cloud strategy and make informed decisions that drive success.

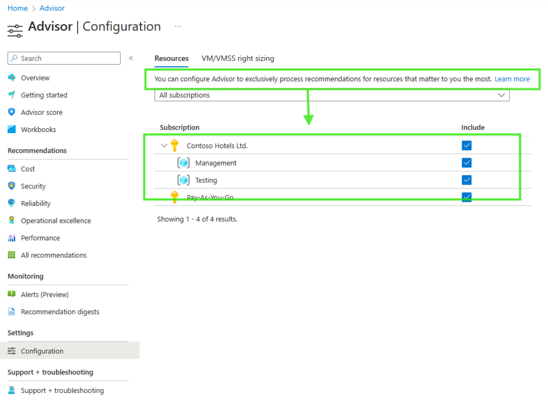

You can configure Advisor to exclusively process recommendations for resources that matter to you the most. Learn more

Subscription	Include
Contoso Hotels Ltd.	☑
Management	☑
Testing	☑
Pay-As-You-Go	☑

Showing 1 - 4 of 4 results.

In this scenario, you need to identify unattached disks in specific resource groups to reduce costs. Azure Advisor is specifically designed for this type of task. By modifying the Advisor configuration to focus on the targeted resource groups, you ensure that:

1. You receive recommendations only for the resource groups you're interested in, reducing noise from irrelevant suggestions.
2. The cost optimization recommendations will include identifying unattached disks, which is one of Advisor's key features.
3. You can easily review and act on these recommendations without needing to create custom queries or implement complex policies.
4. The process is efficient and uses Azure's built-in capabilities, requiring minimal additional setup or maintenance.

Hence, the correct answer is: **Modify the Advisor configuration in Azure Advisor to include only the targeted resource groups, then review its cost optimization recommendations.**

The option that says: **Assign the Billing Administrator role to monitor unattached disks in selected resource groups for potential deletion** is incorrect because the Billing Administrator role is primarily focused on managing billing and subscription-related tasks, not on monitoring or managing Azure resources like unattached disks. This role does not have the permissions necessary to interact with or manage Azure resources, as its scope is limited to billing-related activities.

The option that says: **Use Azure Monitor to create a custom log query that identifies unattached disks in the specified resource groups** is incorrect. Although Azure Monitor can be used to create custom queries, it's not the most straightforward or recommended method for identifying unattached disks. This approach would require significant custom development and may not be as comprehensive as built-in Azure services which are primarily designed for cost optimization.

The option that says: **View the Advisor recommendations under Azure Cost Management and apply filters to focus on the resource groups in question** is incorrect. Azure Cost Management does provide cost analysis and is able to integrate with Azure Advisor. However, for the specific task of identifying unattached disks in particular resource groups, directly configuring Azure Advisor is more appropriate.

References:

https://learn.microsoft.com/en-us/azure/advisor/advisor-overview

https://learn.microsoft.com/en-us/azure/advisor/view-recommendations

Check out this Azure Advisor Cheat Sheet:

https://tutorialsdojo.com/azure-advisor/

42. QUESTION

You have an Azure subscription that contains the resources shown in the following table.

Name	Service	Region
TD1	Virtual machine	Singapore
TD2	Virtual machine	Tokyo
TDVnet1	Virtual network	Singapore
TDVnet2	Virtual network	Tokyo

You provisioned an Azure Bastion named `TD3` in the Singapore region with `TDVnet1` as its virtual network.

For each of the following items, choose Yes if the statement is true or choose No if the statement is false. Take note that each correct item is worth one point.

Questions	Yes	No
`TD3` can only connect to `TD1` using the public IP address of `TD1` .	○	●
`TD3` can connect to `TD1` using the private IP address of `TD1` .	●	○
`TD3` can connect to `TD2` .	●	○

Incorrect

2 / 3 Points

Azure Bastion is a service you deploy that lets you connect to a virtual machine using your browser and the Azure portal. The Azure Bastion service is a fully platform-managed PaaS service that you provision inside your virtual network. It provides secure and seamless RDP/SSH connectivity to your virtual machines directly from the Azure portal over TLS. When you connect via Azure Bastion, your virtual machines don't need a public IP address, agent, or special client software.

Bastion provides secure RDP and SSH connectivity to all of the VMs in the virtual network in which it is provisioned. Using Azure Bastion protects your virtual machines from exposing RDP/SSH ports to the outside world while still providing secure access using RDP/SSH.

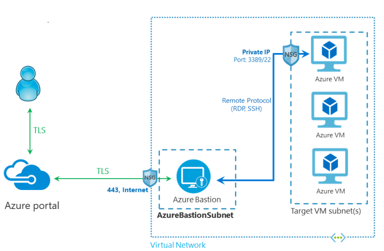

Azure Bastion opens the RDP/SSH connection to your Azure VM by using the private IP address on your VM. You don't need a public IP address on your virtual machine.

Azure Bastion and VNet peering can be used together so that `TD2` and `TD3` can communicate with each other. When VNet peering is configured, you don't have to deploy Azure Bastion in each peered VNet. This means if you have an Azure Bastion host configured in one virtual network (VNet), it can be used to connect to VMs deployed in a peered VNet without deploying an additional bastion host.

You also don't need to apply any NSGs to the Azure Bastion subnet. Since `TD3` and `TD1` are using the same virtual network, they can communicate with each other using their respective private IP addresses. This removes the hassle of managing NSGs each time you need to connect securely to your virtual machines.

Normally virtual machines can connect to other virtual machines located in different virtual networks by their public IP addresses provided that their network security groups permit them.

Hence, this statement is correct: **`TD3` can connect to `TD1` using the private IP address of `TD1` .**

The statement that says: **`TD3` can connect to `TD2`** is incorrect because since `TD3` and `TD2` are in different virtual networks, they will not be able to communicate with each other. For them to be able to connect, their respective virtual networks must have VNet peering or VPN gateway enabled.

The statement that says: **`TD3` can only connect to `TD1` using the public address of `TD1`** is incorrect because `TD3` can connect to `TD1` using the public and private IP addresses of `TD1` . It is highly advisable though to use the private IP address of `TD1` for security reasons.

References:

https://docs.microsoft.com/en-us/azure/bastion/bastion-overview

https://docs.microsoft.com/en-us/azure/bastion/vnet-peering

45. QUESTION

Your company hosts its business-critical Azure virtual machines in the `Australia East` region.

The servers are then replicated to a secondary region using Azure site recovery for disaster recovery.

The `Australia East` region is experiencing an outage and you need to failover to your secondary region.

Which actions should you perform? (Select THREE.)

☐ Reprotect virtual machine.

☐ Run a failback.

☑ Verify if the virtual machines are protected and healthy.

☑ Initiate replication.

☑ Run a test failover.

☐ Run a failover.

Incorrect

A **Recovery Services vault** is a storage entity in Azure that houses data. The data is typically copies of data or configuration information for virtual machines (VMs), workloads, servers, or workstations. You can use Recovery Services vaults to hold backup data for various Azure services such as IaaS VMs (Linux or Windows) and Azure SQL databases.

Recovery Services vaults support System Center DPM, Windows Server, Azure Backup Server, and more. Recovery Services vaults make it easy to organize your backup data while minimizing management overhead.

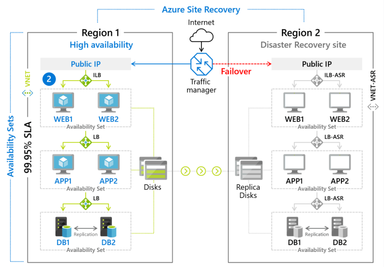

When you enable replication for a VM to set up disaster recovery, the Site Recovery Mobility service extension installs on the VM and registers it with Azure Site Recovery.

During replication, VM disk writes are sent to a cache storage account in the source region. Data is sent from there to the target region, and recovery points are generated from the data. When you fail over a VM during disaster recovery, a recovery point is used to restore the VM in the target region.

To perform a failover, you should complete the following steps:

1. **Verify the VM settings** – Check if the VM is healthy and protected. You also need to verify if the VM is running a support Windows or Linux operation system and if the VM complies with compute, storage and networking requirements.
2. **Run a failover** – In the failover tab, you are required to choose a recovery point. The Azure VM in the target region is created using data from this recovery point.
3. **Reprotect the VM –** After failover, you reprotect the VM in the secondary region so that it replicates back to the primary region.

Hence, the correct answers are:

– Verify if the virtual machines are protected and healthy.

– Run a failover.

– Reprotect the VM.

Initiate replication is incorrect because this is the first step in setting up a disaster recovery for virtual machines. The question states that the servers are already replicated to the secondary region which indicates that it is ready for a failover

Run a failback is incorrect because this option allows you to failback to your primary region and is only executed once the primary region is running as normal again.

Run a test failover is incorrect because you only run a test failover to check if an actual failover will work. This is done during disaster recovery drills.

References:

https://docs.microsoft.com/en-us/azure/site-recovery/site-recovery-overview

https://docs.microsoft.com/en-us/azure/site-recovery/azure-to-azure-tutorial-enable-replication

Check out this Azure Virtual Machines Cheat Sheet:

https://tutorialsdojo.com/azure-virtual-machines/

47. QUESTION

Note: This item is part of a series of questions with the exact same scenario but with a different proposed answer. Each one in the series has a unique solution that may, or may not, comply with the requirements specified in the scenario.

A company has an Azure subscription that contains the following Azure virtual machines:

Name	Status	Virtual network/ Subnet	Public IP SKU
TD1	Running	TDVnet1/TDSub1	None
TD2	Stopped	TDVnet1/TDSub2	Standard
TD3	Running	TDVnet1/TDSub1	Standard
TD4	Stopped	TDVnet1/TDSub2	Basic

You created an Azure public load balancer with the following settings:

- Name: Manila
- SKU: Standard
- Virtual Network: TDVnet1

You need to be able to attach the virtual machines above to the backend pool of Manila.

Solution: Attach a public IP address with an SKU of Basic to TD1

Does this meet the goal?

○ No

● Yes

Incorrect

A **public load balancer** can provide outbound connections for virtual machines (VMs) inside your virtual network. These connections are accomplished by translating their private IP addresses to public IP addresses. Public Load Balancers are used to load balance Internet traffic to your VMs.

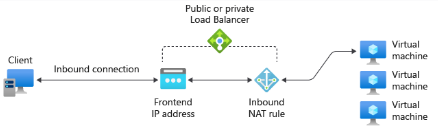

A public IP associated with a load balancer serves as an Internet-facing frontend IP configuration. The Frontend is used to access resources in the backend pool. The frontend IP can be used for members of the backend pool to egress to the Internet.

Remember the following concepts when attaching virtual machines to the backend pool of a load balancer:

– The SKU of the IP address of a virtual machine must match the SKU of the load balancer.

– You can attach virtual machines to the backend pool of a load balancer that does not contain an IP address.

– You can attach virtual machines to the backend pool of a load balancer even if they are in a stopped state.

Without implementing any changes, you would only be able to attach `TD1,` `TD2,` and `TD3` to Manila.

The solution is incorrect because even if you attach a public IP address with an SKU of basic to TD1, you still won't be able to attach it because its SKU must match the SKU of `Manila` which is Standard.

Hence, the correct answer is: **No.**

References:

https://docs.microsoft.com/en-us/azure/virtual-network/ip-services/public-ip-addresses

https://docs.microsoft.com/en-us/azure/virtual-network/ip-services/configure-public-ip-load-balancer

Check out this Azure Load Balancer Cheat Sheet:

https://tutorialsdojo.com/azure-load-balancer/

49. QUESTION

Note: This item is part of a series of questions with the exact same scenario but with a different proposed answer. Each one in the series has a unique solution that may, or may not, comply with the requirements specified in the scenario.

A company has an Azure subscription that contains the following Azure virtual machines:

Name	Status	Virtual network/ Subnet	Public IP SKU
TD1	Running	TDVnet1/TDSub1	None
TD2	Stopped	TDVnet1/TDSub2	Standard
TD3	Running	TDVnet1/TDSub1	Standard
TD4	Stopped	TDVnet1/TDSub2	Basic

You created an Azure public load balancer with the following settings:

- Name: Manila
- SKU: Standard
- Virtual Network: TDVnet1

You need to be able to attach the virtual machines above to the backend pool of Manila.

Solution: Remove the public IP address attached to TD4.

Does this meet the goal?

○ Yes

● No

Incorrect

A **public load balancer** can provide outbound connections for virtual machines (VMs) inside your virtual network. These connections are accomplished by translating their private IP addresses to public IP addresses. Public Load Balancers are used to load balance Internet traffic to your VMs.

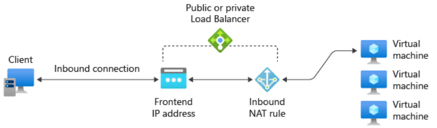

A public IP associated with a load balancer serves as an Internet-facing frontend IP configuration. The Frontend is used to access resources in the backend pool. The frontend IP can be used for members of the backend pool to egress to the Internet.

Remember the following concepts when attaching virtual machines to the backend pool of a load balancer:

– The SKU of the IP address of a virtual machine must match the SKU of the load balancer.

– You can attach virtual machines to the backend pool of a load balancer that does not contain an IP address.

– You can attach virtual machines to the backend pool of a load balancer even if they are in a stopped state.

Without implementing any changes, you would only be able to attach `TD1,` `TD2,` and `TD3` to Manila.

The solution is correct because you can now attach `TD4.` Take note that virtual machines attached to a backend pool do not need a public IP address, even if it will be attached to a public load balancer. Another correct solution is by upgrading the SKU the public IP address of `TD4` to standard.

Hence, the correct answer is: **Yes.**

References:

https://docs.microsoft.com/en-us/azure/virtual-network/ip-services/public-ip-addresses

https://docs.microsoft.com/en-us/azure/virtual-network/ip-services/configure-public-ip-load-balancer

Check out this Azure Load Balancer Cheat Sheet:

https://tutorialsdojo.com/azure-load-balancer/

53. QUESTION

Note: This item is part of a series of case study questions with the exact same scenario but has a different technical requirement. Each one in the series has a unique solution that may or may not comply with the requirements specified in the scenario.

Overview

TechWare Innovations is a technology company with global offices. It collaborates with various partners to deliver its innovative products. Their product designs are maintained as digital files that are critical to their operations.

Existing Environment

TechWare currently operates several types of servers for their business needs, including:

- File servers
- Domain controllers
- Microsoft SQL Server servers

Your network consists of an Azure Active Directory forest called <code>techware.com</code>, with all servers and client computers joined to this domain. Additionally, you have a publicly available application named TechApp, which is composed of three tiers:

- A SQL database
- A presentation tier
- An application logic tier

Each tier consists of ten virtual machines and users exclusively access the web front end via HTTPS.

Planned Changes
TechWare currently operates several types of servers for their business needs, Inc. intends to implement the following infrastructure changes:

- Migrate all tiers of TechApp to Azure.
- Transfer the existing design documents to Azure Blob storage.
- Establish a hybrid directory setup in preparation for a Microsoft 365 migration

Technical Requirements
TechWare must adhere to the following technical requirements:

- Migrate all VMs for TechApp to Azure.
- Minimize the number of open ports between the different components of TechApp.
- Ensure all virtual machines for TechApp are backed up.
- Transfer the design documents to Azure via the Internet.
- Store the design documents in the archive storage tier.
- Secure and provide temporary access for partners to the design documents.
- Prevent the storage of user passwords or password hashes in Azure.
- Utilized unmanaged standard storage for the virtual machines' hard disks.
- Ensure users verify their identity with a mobile phone when joining devices to Microsoft Entra ID.
- Minimize administrative overhead wherever possible.

User Requirements
TechWare has the following user requirements:

- Only users in a group named TechAppGroup can join devices to Entra ID.
- Assign the Owner role at the subscription level to a new user named TechAdmin so that the user can manage all resources, including access control.
- TechAdmin should receive email alerts for service outages.
- The Azure Subscription must allow a new user named NetworkUser to create network objects.

To meet the user requirement for Admin Manager, what actions should you take?

○ Configure the OAuth 2.0 authentication endpoint for app-based access to Azure resources.

● Navigate to the Microsoft Entra Admin Center and configure the group settings.

○ Use Microsoft Exchange Online to create a distribution group for administrative access notifications.

○ Open the Subscriptions blade, select the relevant subscription, and update the Access Control (IAM) permissions.

Incorrect

Configuring the Access Control (IAM) settings is crucial because it allows for the proper assignment of roles and permissions within the Azure environment. Admin Manager needs to be assigned a role such as "Owner" or "Contributor" to gain the appropriate access level for managing the subscription. This ensures that the Admin Manager has the necessary administrative capabilities and can receive email notifications about service outages, enabling them to take timely action.

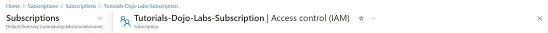
Home > Subscriptions > Subscriptions > Tutorials-Dojo-Labs-Subscription
Subscriptions
Default Directory (tutorialsdojolabstutorialsd.onmi...
Tutorials-Dojo-Labs-Subscription | Access control (IAM)
Subscription

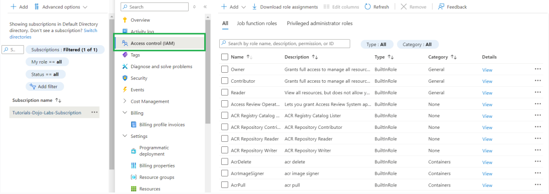

Showing subscriptions in Default Directory directory. Don't see a subscription? Switch directories

Subscriptions : **Filtered (1 of 1)**

My role == **all**

Status == **all**

Subscription name ↑↓

Tutorials-Dojo-Labs-Subscription

Type : **All**

Category : **All**

☐ Name ↑↓	Description ↑↓	Type ↑↓	Category ↑↓	Details
☐ Owner	Grants full access to manage all resourc...	BuiltInRole	General	View
☐ Contributor	Grants full access to manage all resourc...	BuiltInRole	General	View
☐ Reader	View all resources, but does not allow y...	BuiltInRole	General	View
☐ Access Review Operat...	Lets you grant Access Review System ap...	BuiltInRole	None	View
☐ ACR Registry Catalog ...	ACR Registry Catalog Lister	BuiltInRole	None	View
☐ ACR Repository Contri...	ACR Repository Contributor	BuiltInRole	None	View
☐ ACR Repository Reader	ACR Repository Reader	BuiltInRole	None	View
☐ ACR Repository Writer	ACR Repository Writer	BuiltInRole	None	View
☐ AcrDelete	acr delete	BuiltInRole	Containers	View
☐ AcrImageSigner	acr image signer	BuiltInRole	Containers	View
☐ AcrPull	acr pull	BuiltInRole	Containers	View

To assign the service administrator role to Admin Manager, navigate to the Azure Subscriptions blade's Access Control (IAM) settings. This section is essential for managing role-based access control, allowing you to assign specific user roles within your Azure subscription. Selecting the subscription in question and modifying the IAM settings allows you to designate Admin Manager as the service administrator. This ensures that the Admin Manager has the appropriate permissions to manage subscription settings and receive notifications about service interruptions.

Hence, the correct answer is: **Open the Subscriptions blade, select the relevant subscription, and update the Access Control (IAM) permissions.**

The option that says: **Configure the OAuth 2.0 authentication endpoint for app-based access to Azure resources** is incorrect because OAuth 2.0 is primarily used to authenticate and authorize applications to access resources on behalf of users or services, typically for securing APIs or enabling third-party applications to interact with Azure resources. It pertains not to role assignments or managing administrative access within Azure Subscriptions.

The option that says: **Go to the Microsoft Entra Admin Center and update the directory properties** is incorrect because updating directory properties in Microsoft Entra ID typically involves configuring settings related to the directory itself, not assigning roles for subscription management.

The option that says: **Use Microsoft Exchange Online to create a distribution group for administrative access notifications** is incorrect because Microsoft Exchange Online is just an email and messaging service. Creating a distribution group can help send notifications to multiple recipients, but it does not assign Azure RBAC roles or grant subscription-level administrative permissions. The requirement is to provide administrative access within the Azure subscription, which must be done through Access Control (IAM).

References:

https://learn.microsoft.com/en-us/azure/role-based-access-control/role-assignments-portal

https://learn.microsoft.com/en-us/azure/role-based-access-control/overview

Check out this Azure Role-Based Access Control (RBAC) Cheat Sheet:

https://tutorialsdojo.com/azure-role-based-access-control-rbac/

55. QUESTION

You are tasked with planning a deployment for a new containerized application environment, `TD-ContAppEnv1`, within a company's Azure subscription that consists of the following virtual networks:

Name	Location	IP Address Space	Subnet Names
TD-VNetA	West US	10.0.0.0/16	TD-SubnetA1
TD-VNetB	West US	172.30.0.0/16	TD-SubnetB1, TD-SubnetB2
TD-VNetC	West US	192.168.100.0/22	TD-SubnetC1

Subnets are configured with the following IP address spaces:

Subnet Name	IP Address Space
TD-SubnetA1	10.0.1.0/24
TD-SubnetB1	172.30.0.0/18
TD-SubnetB2	172.30.64.0/18
TD-SubnetC1	192.168.100.0/24

The deployment requirements for `TD-ContAppEnv1` are as follows:

- Workload profiles Environment type
- An isolated subnet.
- The subnet must be smallest possible subnet.

Given the requirements for `TD-ContAppEnv1`, to which virtual networks can this environment be connected, and what is the smallest possible subnet mask that should be used for the new subnet?

Select the correct answer from the drop-down list of options. Each correct selection is worth one point.

Virtual Network: *TD-VNetA, TD-VNetB, or TD-VNetC*

Subnet Mask: */24* ***(/26)***

Incorrect 1 / 2 Points

Azure Container Apps is a serverless platform that allows you to maintain less infrastructure and save costs while running containerized applications. Instead of worrying about server configuration, container orchestration, and deployment details, Container Apps provides all the up-to-date server resources required to keep your applications stable and secure.

Azure Container Apps run in the context of an environment with its own virtual network (VNet).

By default, your Container App environment is created with a VNet that is automatically generated for you. You can provide an existing VNet when you create an environment for fine-grained control over your network. The network type can't be changed once you create an environment with either a generated or existing VNet.

Generated VNets take on the following characteristics.

They are:

– inaccessible to you as they're created in Microsoft's tenant

– publicly accessible over the internet

– only able to reach internet-accessible endpoints

Further, they only support a limited subset of networking capabilities, such as ingress IP restrictions and container app-level ingress controls.

Use an existing VNet if you need more Azure networking features, such as:

– Integration with Application Gateway

– Network Security Groups

– Communication with resources behind private endpoints in your virtual network

The available VNet features depend on your environment selection.

Container Apps has two different environment types, which share many of the same networking characteristics with some key differences.

Workload profiles – Supports user-defined routes (UDR) and egress through NAT Gateway. The minimum required subnet size is /27.

Consumption only – Doesn't support user-defined routes (UDR), egress through NAT Gateway, peering through a remote gateway, or other custom egress. The minimum required subnet size is /23.

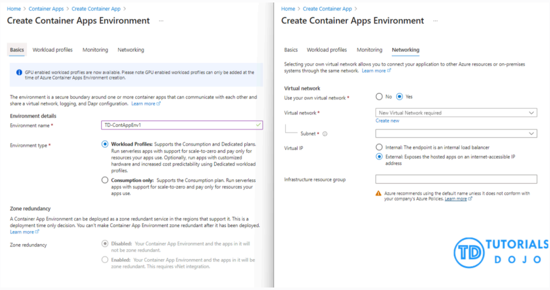

Create Container Apps Environment

GPU enabled workload profiles are now available. Please note GPU enabled workload profiles can only be added at the time of Azure Container Apps Environment creation.

The environment is a secure boundary around one or more container apps that can communicate with each other and share a virtual network, logging, and Dapr configuration. Learn more

Environment details

Environment name * TD-ContAppEnv1

Environment type *

● **Workload Profiles:** Supports the Consumption and Dedicated plans. Run serverless apps with support for scale-to-zero and pay only for resources your apps use. Optionally, run apps with customized hardware and increased cost predictability using Dedicated workload profiles.

○ **Consumption only:** Supports the Consumption plan. Run serverless apps with support for scale-to-zero and pay only for resources your apps use.

Zone redundancy

A Container App Environment can be deployed as a zone redundant service in the regions that support it. This is a deployment time only decision. You can't make Container App Environment zone redundant after it has been deployed. Learn more

Zone redundancy

● **Disabled:** Your Container App Environment and the apps in it will not be zone redundant.

○ **Enabled:** Your Container App Environment and the apps in it will be zone redundant. This requires vNet integration.

Create Container Apps Environment

Selecting your own virtual network allows you to connect your application to other Azure resources or on-premises systems through the same network. Learn more

Virtual network

Use your own virtual network * ○ No ● Yes

Virtual network * New Virtual Network required

Create new

Subnet *

Virtual IP

○ Internal: The endpoint is an internal load balancer

● External: Exposes the hosted apps on an internet-accessible IP address

Infrastructure resource group

Azure recommends using the default name unless it does not conform with your company's Azure Policies. Learn more

The deployment requirements state that the environment needs an isolated subnet, and all three virtual networks (TD-VNetA, TD-VNetB, and TD-VNetC) have available subnets that can accommodate the new isolated subnet for `TD-ContAppEnv1`. It also requires workload profiles environment type, meaning the minimum required subnet size is /27. However, since /27 is not in the given options, the next smallest possible subnet mask that can accommodate a new subnet is /26.

Therefore, the correct answers are:

Virtual Network: **TD-VNetA, TD-VNetB, or TD-VNetC**

Subnet Mask: **/26**

References:

https://learn.microsoft.com/en-us/azure/container-apps/networking?tabs=workload-profiles-env%2Cazure-cli

https://learn.microsoft.com/en-us/azure/container-apps/overview

Check out this Azure Virtual Network Cheat Sheet:

https://tutorialsdojo.com/azure-virtual-network-vnet/